

# An Actor-Critic Algorithm for Sequence Prediction

Dzmitry Bahdanau, Philemon Brakel, Kelvin Xu, Ryan Lowe, Joelle Pineau, Aaron Courville, Yoshua Bengio

Presented By Payam Nikdel


# RL Background

- Agent learning sequence of actions(  $a_t$  )
- Observes outcomes (state  $s_{t+1}$ , rewards  $r_{t+1}$  ) of those actions
- **Goal**: find a Policy  $\pi$  that maximizes the return  $R$

$$\pi = p(a|s)$$

$$R = \sum_{t=0}^T \gamma^t r_{t+1} \quad \gamma: \text{Discount rate}$$

$$V(s_t): \text{Value function} = E_{a \sim \pi}[R | S_t]$$



# Policy Evaluation

Method for calculating the value function for a policy

- **Monte Carlo Learning:**

Wait until end of episode to observe R


$$V(s_t) = V(s_t) + \alpha [R - V(s_t)]$$

- **TD(0) Learning:**

bootstrap off previous estimate of v


$$V(s_t) = V(s_t) + \alpha \underbrace{\left[ (r_t + \gamma V(s_{t+1})) - V(s_t) \right]}_{\text{TD-error}}$$

## Monte Carlo



$$V(s_t) = V(s_t) + \alpha [R - V(s_t)]$$

## TD(0)



$$V(s_t) = V(s_t) + \alpha [ (r_t + \gamma V(s_{t+1})) - V(s_t) ]$$

# TD learning, adapted from Sutton and Barto (2017)

**Input:** the policy  $\pi$  to be evaluated

**Output:** value function  $V$

initialize  $V$  arbitrarily, e.g., to 0 for all states

**for each episode do:**

    initialize state  $s$

**for each step of episode, state  $s$  is not terminal do:**

$a \leftarrow$  action given by  $\pi$  for  $s$

        take action  $a$ , observe  $r, s'$

$V(s) \leftarrow V(s) + \alpha[r + \gamma V(s') - V(s)]$

$s \leftarrow s'$

**end**

**end**

# Actor-Critic



Initially: act out some policy



# Actor-Critic



Initially: act out some policy



Critic  $\equiv$  policy evaluation



# Actor-Critic



Actor  $\equiv$  policy improvement




Critic  $\equiv$  policy evaluation



# Actor-Critic


- Have a parametrized value function  $V$  (the critic) and policy  $\pi$  (the actor)
- Actor takes actions according to  $\pi$ , critic 'criticizes' them with TD error
- TD error drives learning of both actor and critic



Sutton & Barto, 1998

# Actor-Critic


- Critic learns with usual TD learning, or with LSTD
- Actor learns according to the policy gradient theorem



Sutton & Barto, 1998

# Actor-Critic for Sequence Prediction

- **Actor:** some function with parameters  $\theta$  that predicts sequence one token at a time (i.e. generates 1 word at a time)
- **Critic:** some function with parameters  $\phi$  that computes the TD-error of decisions made by actor, which is used for learning



# Why Actor-Critic?

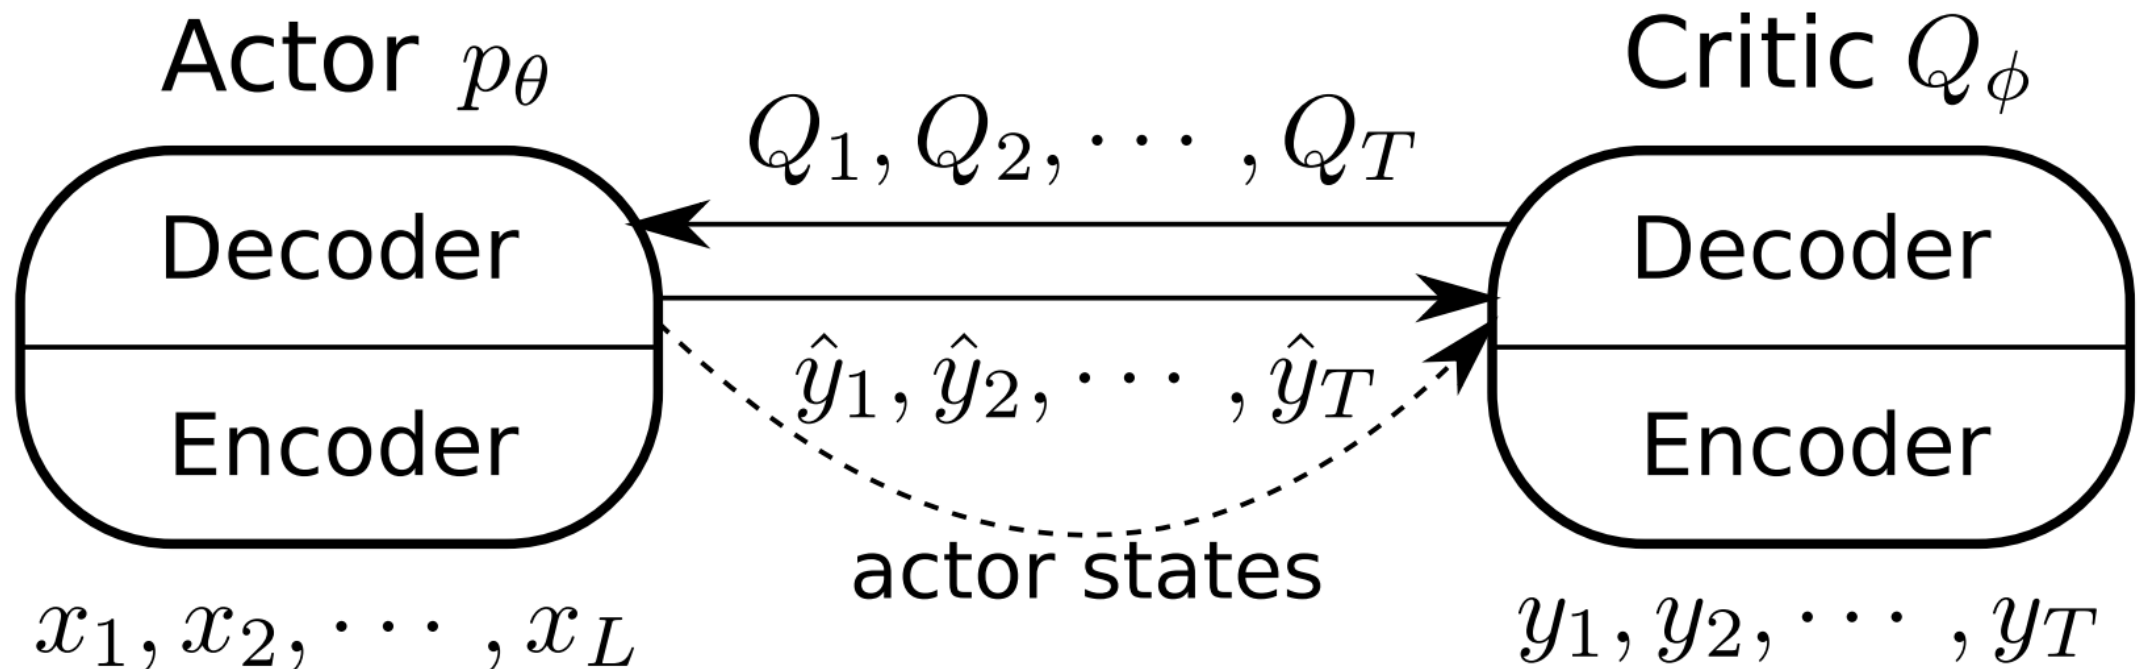
- **Teacher forcing** (like maximum log-likelihood)<sup>[1]</sup>:
  - During training, tokens come from training data
  - At generation time, tokens are previous generated y
    - leads to **discrepancies between train and test time.**
  - actor-critic, can **condition on actor's previous outputs**
- Allows for the direct **optimization of a task-specific score**, e.g. BLEU, rather than log-likelihood

[1]Maximum likelihood training can be suboptimal. ( Bengio et al 2015)

# Actor & Critic Notation

- Critic:
  - Separate RNN parametrized by  $\phi$
  - Conditioned on **actor generated sequence**  $\hat{Y}_{1,\dots,t} = (\hat{y}_1, \dots, \hat{y}_t)$  and **ground-truth output**  $Y$
  - Produces the estimates  $Q^{\pi}(a; \hat{Y}_{1,\dots,t})$  for all  $a \in A$
  - *Critic Q-values are only required during training! No critic in test time*
- Actor:
  - Actor policy  $p(a; \hat{Y}_{1,\dots,t}, X)$  is conditioned on **outputs so far**  $\hat{Y}_{1,\dots,t}$  and the **input**  $X$

# Actor & Critic Notation





# Algorithm: Policy gradients for sequence prediction

- Denote  $V$  as the expected reward under  $\pi_\theta$

**Proposition 1** *The gradient  $\frac{dV}{d\theta}$  can be expressed*

*using  $Q$  values of intermediate actions:*

$$\frac{dV}{d\theta} = \mathbb{E}_{\hat{Y} \sim p(\hat{Y})} \sum_{t=1}^T \sum_{a \in \mathcal{A}} \frac{dp(a|\hat{Y}_{1\dots t-1})}{d\theta} Q(a; \hat{Y}_{1\dots t-1})$$

# Algorithm: Policy gradients for sequence prediction

- 2: **while** Not Converged **do**
- 3:   Receive a random example  $(X, Y)$ .
- 4:   Generate a sequence of actions  $\hat{Y}$  from  $p'$ .
- 5:   Compute targets for the critic

$$q_t = r_t(\hat{y}_t; \hat{Y}_{1..t-1}, Y) + \sum_{a \in \mathcal{A}} p'(a | \hat{Y}_{1..t}, X) \hat{Q}'(a; \hat{Y}_{1..t}, Y)$$

# Algorithm: Policy gradients for sequence prediction

6: Update the critic weights  $\phi$  using the gradient

$$\frac{d}{d\phi} \left( \sum_{t=1}^T \left( \hat{Q}(\hat{y}_t; \hat{Y}_{1\dots t-1}, Y) - q_t \right)^2 + \lambda C \right)$$

# Algorithm: Policy gradients for sequence prediction

- 7: Update actor weights  $\theta$  using the following gradient estimate

$$\frac{dV(X, Y)}{d\theta} = \sum_{t=1}^T \sum_{a \in \mathcal{A}} \frac{dp(a | \hat{Y}_{1 \dots t-1}, X)}{d\theta} \hat{Q}(a; \hat{Y}_{1 \dots t-1}, Y)$$

# Deep implementation: ACTOR

- use an RNN with ‘soft-attention’ (Bahdanau et al., 2015)
- Encode source sentence  $X$  with bi-directional GRU
- Compute **weighted sum over  $x$ 's** at each time step using weights  $\alpha$

$$c_i = \sum_{j=1}^{T_x} \alpha_{ij} h_j.$$

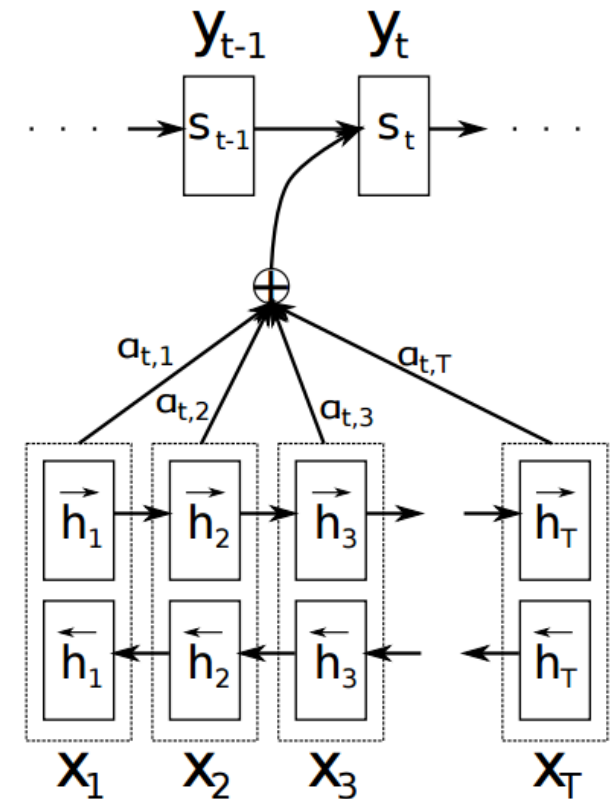


Figure 1: The graphical illustration of the proposed model trying to generate the  $t$ -th target word  $y_t$  given a source sentence  $(x_1, x_2, \dots, x_T)$ .

# Deep implementation: CRITIC

- use the same architecture, except conditioned on  $Y$  instead of  $X$
- Input: the sequence generated so far  $\hat{Y}_{1\dots t}$ , and the ground-truth sequence  $Y$
- Output: Q-value prediction

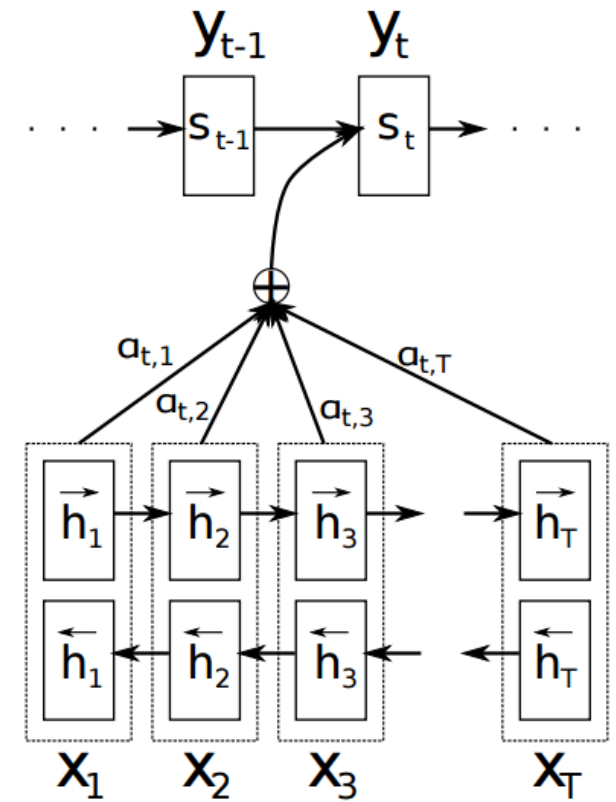


Figure 1: The graphical illustration of the proposed model trying to generate the  $t$ -th target word  $y_t$  given a source sentence  $(x_1, x_2, \dots, x_T)$ .

# Tricks: Target network

- Similarly to DQN, use a **target network**
- In particular, have both **delayed actor  $\mathbf{p}'$**  and a **delayed critic  $\mathbf{Q}'$** , with params  $\theta'$  and  $\phi'$ , respectively
- Use this delayed values to compute target for critic:

$$q_t = r_t(\hat{y}_t; \hat{Y}_{1...t-1}, Y) + \sum_{a \in \mathcal{A}} p'(a | \hat{Y}_{1...t}, X) \hat{Q}'(a; \hat{Y}_{1...t}, Y)$$



# Tricks: Target network

- After updating actor and critic, update delayed actor and critic using a linear interpolation:

8: Update delayed actor and target critic, with a constant  $\tau \ll 1$ :

$$\theta' = \tau\theta + (1 - \tau)\theta'$$

$$\phi' = \tau\phi + (1 - \tau)\phi'$$

# Tricks: variance penalty

- Problem: critic can have **high variance** for words that are **rarely sampled**
- Solution: artificially reduce values of rare actions by introducing a **variance regularization** term:

$$C = \sum_a \left( \hat{Q}(a; \hat{Y}_{1...t-1}) - \frac{1}{|\mathcal{A}|} \sum_b \hat{Q}(b; \hat{Y}_{1...t-1}) \right)^2,$$

# Tricks: Reward decomposition

- Could train critic using all the score at the last step, but this signal is sparse
- Want to improve learning of critic (and thus the actor) by **providing rewards at each time step**
- If final reward is  $R(\hat{Y})$ , decompose the reward into scores for all prefixes:  $(R(\hat{Y}_{1...1}), R(\hat{Y}_{1...2}), \dots, R(\hat{Y}_{1...T}))$
- Then the reward at time step  $t$  is:

$$r_t(\hat{y}_t) = R(\hat{y}_{1...t}) - R(\hat{y}_{1...t-1})$$

# Tricks: Pre-training

- It will take forever to learn, if we start off with a random actor and critic
- **Use pre-training:** first train actor to increase likelihood of correct answer (Maximizing  $\log p(y_{t+1} | Y_{1...t}, X)$ )
- Then, train critic by feeding samples from a fixed actor

# Experiments

## 1. Tested on a synthetic spelling correction task

- Dataset generated by randomly replacing a character using a random character
- One billion word dataset (no chance for overfitting)
- Used a character error rate (CER) as reward
- models:
  - Maximum likelihood
  - Actor-critic
  - Two version of REINFORCE
    - I. Exactly as in [1]
    - II. Use critic as a base line for REINFORCE

[1] SEQUENCE LEVEL TRAINING WITH RECURRENT NEURAL NETWORKS (Ranzato et al 2015)




Figure 2: Progress of log-likelihood (LL), REINFORCE (RF) and actor-critic (AC) training in terms of BLEU score on the training (train) and validation (valid) datasets. LL\* stands for the annealing phase of log-likelihood training. The curves start from the epoch of log-likelihood pretraining from which the parameters were initialized.

Table 1: Character error rate of different methods on the spelling correction task. In the table  $L$  is the length of input strings,  $\eta$  is the probability of replacing a character with a random one. LL stands for the log-likelihood training, AC and RF-C and for the actor-critic and the REINFORCE-critic respectively, AC+LL and RF-C+LL for the combinations of AC and RF-C with LL.

Setup	Character Error Rate				
	LL	AC	RF-C	AC+LL	RF-C+LL
$L = 10, \eta = 0.3$	17.81	17.24	17.82	<b>16.65</b>	16.97
$L = 30, \eta = 0.3$	18.4	17.31	18.16	<b>17.1</b>	17.47
$L = 10, \eta = 0.5$	38.12	35.89	35.84	<b>34.6</b>	35
$L = 30, \eta = 0.5$	40.87	37.0	37.6	<b>36.36</b>	36.6



# Experiments

## 2. Tested on a real-world German to English and English to French machine translation

### 1. IWSLT 2014 :

- 153,000 aligned sentence pairs in training set (German to English )
- Changed the actor to 256 hidden units instead of GRU to compare it with REINFORCE [1] model.

### 2. WMT 14 :

- More than 12 million examples

[1] SEQUENCE LEVEL TRAINING WITH RECURRENT NEURAL NETWORKS (Ranzato et al 2015)

Table 3: Our IWSLT 2014 machine translation results with a bidirectional recurrent encoder compared to the previous work. Please see Table 1 for an explanation of abbreviations. The asterisk identifies results from (Wiseman & Rush, 2016).

Decoding method	Model						
	LL*	BSO*	LL	RF-C	RF-C+LL	AC	AC+LL
greedy search	22.53	23.83	25.82	27.42	<b>27.7</b>	27.27	27.49
beam search	23.87	25.48	27.56	27.75	28.3	27.75	<b>28.53</b>

Table 4: Our WMT 14 machine translation results compared to the previous work. Please see Table 1 for an explanation of abbreviations. The apostrophy and the asterisk identify results from (Bahdanau et al., 2015) and (Shen et al., 2015) respectively.

Decoding method	Model					
	LL'	LL*	MRT *	LL	AC+LL	RF-C+LL
greedy search	n/a	n/a	n/a	29.33	<b>30.85</b>	29.83
beam search	28.45	29.88	<b>31.3</b>	30.71	31.13	30.37

Word	Words with largest $\hat{Q}$
one	and(6.623) there(6.200) but(5.967)
of	that(6.197) one(5.668) &apos;s(5.467)
them	that(5.408) one(5.118) i(5.002)
i	that(4.796) i(4.629) ,(4.139)
want	want(5.008) i(4.160) &apos;t(3.361)
to	to(4.729) want(3.497) going(3.396)
tell	talk(3.717) you(2.407) to(2.133)
you	about(1.209) that(0.989) talk(0.924)
about	about(0.706) .(0.660) right(0.653)
here	.(0.498) ?(0.291) -(0.285)
.	.(0.195) there(0.175) know(0.087)
$\emptyset$	.(0.168) $\emptyset$ (-0.093) ?(-0.173)

**Table 3:** The best 3 words according to the critic at intermediate steps of generating a translation. The numbers in parentheses are the value predictions  $\hat{Q}$ . The German original is “über eine davon will ich hier erzählen .” The reference translation is “and there’s one I want to talk about”.

